

CS 234: Assignment #2

Due date: 5/5 11:59 PM PST

These questions require thought, but do not require long answers. Please be as concise as possible.

We encourage students to discuss in groups for assignments. **However, each student must finish the problem set and programming assignment individually, and must turn in her/his assignment.** We ask that you abide by the university Honor Code and that of the Computer Science department, and make sure that all of your submitted work is done by yourself. If you have discussed the problems with others, please include a statement saying who you discussed problems with. Failure to follow these instructions will be reported to the Office of Community Standards. We reserve the right to run a fraud-detection software on your code.

Please review any additional instructions posted on the assignment page at <http://cs234.stanford.edu/assignment2>. When you are ready to submit, please follow the instructions on the course website. **Make sure you test your code using the provided commands and do not edit outside of the marked areas.**

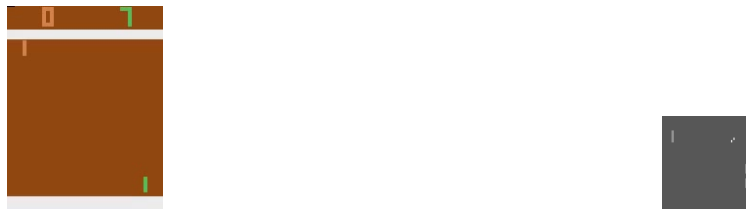
You'll need to download the starter code and fill the appropriate functions following the instructions from the handout and the code's documentation. Training DeepMind's network on Pong takes roughly **12 hours on GPU**, so **please start early!** We will give you access to an Azure GPU cluster. You'll find the setup instructions on the course assignment page.

1 Introduction

In this assignment we will implement **deep Q learning**, following DeepMind's paper ([2] and [3]) that learns to play Atari from raw pixels. The purpose is to understand the effectiveness of deep neural network as well as some of the techniques used in practice to stabilize training and achieve better performance. You'll also have to get comfortable with `Tensorflow`. We will train our networks on the `Pong-v0` environment from OpenAI gym, but the code can easily be applied to any other environment.

The Atari environment from OpenAI gym returns observations of size $(210 \times 160 \times 3)$, the last dimension corresponding to the RGB channels filled with values between 0 and 255 (`uint8`). To reduce the dimension of the input image, we will apply some preprocessing to the observation:

- each time we execute an action, we keep doing it for 4 time steps. We return a pixel-wise max-pooling of the 4 consecutive frames. This reduces the frequency of decisions without impacting the performance too much and enables us to play 4 times more games while training.
- crop the image, convert it to grey scale



(a) Original input ($210 \times 160 \times 3$) with RGB colors (b) After preprocessing in grey scale of shape ($80 \times 80 \times 1$)

Figure 1: Pong-v0 environment

In Pong, one player wins if the ball passes through the other player. Winning a game gives a reward of 1, while losing gives a negative reward of -1. An episode is over when one of the two players reaches 21 wins. Thus, the final score is between -21 (lost all) or +21 (won all). Our agent plays against a decent hard-coded AI player. Human performance is -3. If you go to the end of the homework successfully, you will train an AI agent with super-human performance, reaching at least +5 (hopefully more!).

2 Q-learning

Tabular setting In the *tabular setting* (see assignment 1), we maintained a table $Q(s, a)$ for each tuple state-action. Given an experience sample (s, a, r, s') , our update rule was

$$Q(s, a) = Q(s, a) + \alpha \left(r + \gamma \max_{a' \in A} Q(s', a') - Q(s, a) \right), \quad (1)$$

where $\alpha \in \mathbb{R}$ is the learning rate, γ the discount factor.

Approximation setting Due to the scale of Atari environments, we cannot reasonably learn and store a Q value for each state-action tuple. We will represent our Q values as a function $Q_\theta(s, a)$ where θ are parameters of the function (typically, neural network's weights). In this *approximation* setting, our update rule becomes

$$\theta = \theta + \alpha \left(r + \gamma \max_{a' \in A} Q_\theta(s', a') - Q_\theta(s, a) \right) \nabla_\theta Q_\theta(s, a). \quad (2)$$

In other words, we are trying to minimize

$$L(\theta) = \mathbf{E}_{s,a,r,s'} \left[r + \gamma \max_{a' \in A} Q_\theta(s', a') - Q_\theta(s, a) \right]^2 \quad (3)$$

Target Network DeepMind's paper [2] [3] maintains two sets of parameters, θ (to compute $Q(s, a)$) and θ^- (target network, to compute $Q(s', a')$) s.t. our update rule becomes

$$\theta = \theta + \alpha \left(r + \gamma \max_{a' \in A} Q_{\theta^-}(s', a') - Q_\theta(s, a) \right) \nabla_\theta Q_\theta(s, a). \quad (4)$$

The target network's parameters are updated with the Q-network's parameters occasionally and are kept fixed between individual updates. Note that when computing the update, we don't compute gradients with respect to θ^- (these are considered fixed weights).

Replay Memory As we play, we store our transitions (s, a, r, s') in a buffer. Old examples are deleted as we store new transitions. To update our parameters, we *sample* a minibatch from the buffer and perform a stochastic gradient descent update.

ϵ -greedy exploration strategy During training, we use an ϵ -greedy strategy. DeepMind's paper [2] [3] decreases ϵ from 1 to 0.1 during the first million steps. At test time, the agent choses a random action with probability $\epsilon_{soft} = 0.05$.

There are several things to be noted:

1. In this assignment, we will update θ every `learning_freq` steps by using a minibatch of experiences sampled from the replay buffer.
2. DeepMind's deep Q network takes as input the state s and outputs a vector of size = number of actions. In the pong environment, we have 6 actions, thus $Q_\theta(s) \in \mathbb{R}^6$.
3. The input of the deep Q network is the concatenation 4 consecutive steps, which results in an input after preprocessing of shape $(80 \times 80 \times 4)$.

3 Setup - TestEnv (15pts)

Before running our code on Atari, it is crucial to test our code on a test environment. You should be able to run your models on CPU in no more than a few minutes on the following environment:

- 3 states: 1, 2, 3
- 4 actions: 1, 2, 3, 4. Action $1 \leq i \leq 3$ goes to state i , while action 4 makes the agent stay in the same state.
- one episode lasts 5 time steps and always starts in state 1.
- rewards depend on history. Going to state i gives a reward $r(i)$ such that $r(1) = -0.1, r(2) = 0, r(3) = 0.1$. If we were in state 2 before the transition, then the reward is multiplied by -10 .




Figure 2: Example of a path in the Test Environment

You can check the code of the environment in `utils/test_env.py`. We'll use the same learning strategy as for Atari.

1. (**written** 2pts) General - Why do we use the last 4 time steps as input to the network?
2. (**written** 2pts) General - What would be a benefit of experience replay? What would be a benefit of the target network?
3. (**written** 2pts) General - What would be a benefit of representing the Q function as $Q(s) \in \mathbb{R}^{\text{num actions}}$ instead of $Q(s, a)$?
4. (**written** 5pts) What's the optimal achievable reward for this environment?
5. (**coding** 4pts) Implement `get_action` and `update` functions in `q1_schedule.py`. Test your implementation by running `python q1_schedule.py`.

4 Linear Approximation (30pts)

1. (**written** 5pts) Show that (1) and (2) are exactly the same when $Q_\theta(s, a) = \theta^T \phi(s, a)$, where $\theta \in \mathbf{R}^{|S||A|}$, $\phi : S \times A \rightarrow \mathbf{R}^{|S||A|}$, and

$$\phi(s, a)_{s', a'} = \begin{cases} 1 & \text{if } s' = s, a' = a \\ 0 & \text{otherwise} \end{cases}$$

2. (**written** 5pts) Derive the gradient with regard to the value function parameter $\theta \in \mathbb{R}^n$ given $Q_\theta(s, a) = \theta^T \phi(s, a)$ for any function $\phi(s, a) \mapsto x \in \mathbb{R}^n$ and write the update rule for θ .
3. (**coding** 15pts) Implement linear approximation with Tensorflow. Tensorflow will be particularly useful for more complex functions, and this question will setup the whole pipeline (this step is crucial for the rest of the assignment). You'll need to fill the relevant functions in `q2_linear.py`. Test your code with `python q2_linear.py` that will run linear approximation with Tensorflow on the test environment.
4. (**written** 5pts) Do you reach the optimal achievable reward on the test env? Attach the plot `scores.png` to your writeup.

Each experiment creates a folder in `results/` and load a config from `configs/`. You can edit the configuration file.

5 Implementing DeepMind's DQN (40pts)

1. (**coding** 10pts) Implement the deep Q-network as described in [2] by filling `get_q_values_op` in `q3_nature.py`. The rest of the code inherits from what you wrote for linear approximation. Be sure that it works! Test your implementation on the test environment by running `python q3_nature.py`.
2. (**written** 5pts) Attach the plot of scores. Compare the model with linear approximation. What about the final performance? Training time?
3. (**written** 5pts) What's the number of parameters of this model if the input to the Q network is a tensor of shape (80, 80, 4)? Compare with linear approximation.
4. (**coding and written** 5pts). Now, we're ready to train on the Atari Pong-v0 environment. First, launch linear approximation on pong with `python q4_train_atari_linear.py`. What do you notice?
5. (**coding and written** 10 pts). In this question, we'll train the agent with DeepMind's architecture. Run `python q5_train_atari_nature.py`. We provide you with a default config file. Feel free to edit it. You can also decide to change the exploration strategy and learning rate schedule, that are both linearly decaying at this point. We expect you to run for at least 5 million steps (default in the config file). We provide you with our plot of performance (evaluation score) to help you debug your network.

Deliverable Include your own plot to the writeup. You should get a score of at least 5 after 5 million time steps (note that according to DeeMind's paper, human's performance is -3). Include your changes to the config to the writeup, if any.

As the training time is roughly **12 hours**, you may want to check after a few epochs that your network is making progress. A good way to check that is

- look at the progress of the evaluation score printed on terminal (it should start at -21 and increase).




Figure 3: Evolution of the evaluation score on Pong-v0 for the first 5 million time steps. One epoch corresponds to 250k time steps.

- the max of the q values should also be increasing
 - the standard deviation of q shouldn't be too small. Otherwise it means that all states have similar q values
 - we advise you to use Tensorboard. The starter code writes summaries of a bunch of useful variables that can help you monitor the training process. More on how to use Tensorboard on the assignment page.
6. (**written** 5pts) Compare the performance with previous linear Q-network. How can you explain the gap in the performance?

6 Bonus: Architecture of the Network (10pts)

This is an open question. Design a different network architecture or try a more recent paper, like Double Q Learning [1] or Dueling Networks [4]. Explain your architecture and implement it in `q6_bonus_question.py`. You can build on the existing code architecture or start from scratch. Note that this is a **bonus** question.

7 Uniform Approximation (15pts)

Recall the universal approximation theorem that states a single-hidden-layer neural network with a linear output unit can approximate any continuous function arbitrarily well, given enough hidden units. In this question, we explore the uniform approximation property of neural networks in the case of boolean functions. Assume the input x is a binary vector of length n , $x \in \{0, 1\}^n$. Let the activation function be the step function: $f(z) = 1$ if $z > 0$ and 0 otherwise. Note that the weights can still be real numbers. The output will be the result of a single unit and hence binary, either 0 or 1. We use such neural networks to implement boolean functions.

- (a) (**written** 3pts) Let $n = 2$ and the input be a pair of binary values. Consider a neural network with just a single output unit and no hidden units, i.e., $y = f(w^T x + b)$ where y is the output. Choose w and b to implement boolean AND (i.e., $y = 1$ only when $x = (1, 1)$). Choose w and b to implement boolean OR.
- (b) (**written** 3pts) Under the same conditions as above, what boolean function of two variables cannot be represented? Briefly explain why the function cannot be implemented.
- (c) (**written** 3pts) Suppose we now allow a single layer of hidden units, i.e., $y = f(w^T z + b)$, $z_j = f(w_j^T x + b_j)$ where z_j 's present components of z . Construct a neural network that represents the boolean function you provided in Part (b). There is no restrictions on the number of hidden units, but try to keep it as simple as possible.
- (d) (**written** 3pts) Describe a general procedure to construct a neural network with single hidden layer to represent any boolean function.
- (e) (**written** 3pts) The result from Part (d) implies that a single hidden layer is sufficient to implement any boolean function. Why would you want to consider neural networks with multiple hidden layers?

References

- [1] Hado van Hasselt, Arthur Guez, and David Silver. “Deep Reinforcement Learning with Double Q-learning”. In: *CoRR* abs/1509.06461 (2015). URL: <http://arxiv.org/abs/1509.06461>.
- [2] Volodymyr Mnih et al. “Human-level control through deep reinforcement learning”. In: *Nature* 518.7540 (2015), pp. 529–533.
- [3] Volodymyr Mnih et al. “Playing Atari With Deep Reinforcement Learning”. In: *NIPS Deep Learning Workshop*. 2013.
- [4] Ziyu Wang, Nando de Freitas, and Marc Lanctot. “Dueling Network Architectures for Deep Reinforcement Learning”. In: *CoRR* abs/1511.06581 (2015). URL: <http://arxiv.org/abs/1511.06581>.